Compact Hardware Implementations of ChaCha, BLAKE, Threefish, and Skein on FPGA

Nuray At, Jean-Luc Beuchat, Eiji Okamoto, İsmail San, and Teppei Yamazaki

Abstract—The cryptographic hash functions BLAKE and Skein are built from the ChaCha stream cipher and the tweakable Threefish block cipher, respectively. Interestingly enough, they are based on the same arithmetic operations, and the same design philosophy allows one to design lightweight coprocessors for hashing and encryption. The key element of our approach is to take advantage of the parallelism of the algorithms considered in this work to deeply pipeline our Arithmetic and Logic Units, and to avoid data dependencies by interleaving independent tasks. We show for instance that a fully autonomous implementation of BLAKE and ChaCha on a Xilinx Virtex-6 device occupies 144 slices and three memory blocks, and achieves competitive throughputs. In order to offer the same features, a coprocessor implementing Skein and Threefish requires a substantial higher slice count.

Index Terms—Ciphers , cryptography, coprocessors, field programmable gate arrays.

I. INTRODUCTION

THE cryptographic hash functions BLAKE [1] and Skein [2] are built from the ChaCha stream cipher [3] and the tweakable Threefish block cipher [2], respectively. It is therefore tempting to design compact unified hardware architectures able to hash and encrypt a message. We extend here the work presented in [4], [5], and propose novel hardware architectures for Threefish decryption, ChaCha, and BLAKE. The key element of our approach is to:

- take advantage of the parallelism of the algorithms to deeply pipeline our Arithmetic and Logic Units (ALUs);
- show that a careful scheduling allows us to interleave independent tasks and avoid pipeline bubbles.

The rest of the article is organized as follows: after a brief overview of Threefish (Section II), Skein (Section III), ChaCha (Section IV), and BLAKE (Section V), we describe our design philosophy and compact hardware implementations (Section VI). We discuss our implementation results on Xilinx Virtex-6 Field-Programmable Gate Arrays (FPGAs) in Section VII and conclude in Section VIII.

TABLE I
NUMBER OF ROUNDS OF THREEFISH FOR DIFFERENT KEY SIZES

Key size [bits]	# 64-bit words N_w	# rounds N_r	Block size N_b [bytes]
256	4	72	32
512	8	72	64
1024	16	80	128

Throughout this article, all operands are w-bit unsigned integers and the following notation is adopted:

- $\boxplus$ and $\boxminus$: addition and subtraction modulo 2^w;
- $\vee$ and $\wedge$: bitwise OR and bitwise AND;
- $\oplus$: bitwise exclusive OR;
- $\ggg \alpha$: rotation by α bits to the right;
- $\lll \alpha$: rotation by α bits to the left.

II. THE THREEFISH BLOCK CIPHER

The design philosophy of Threefish is that "a larger number of simple rounds is more secure than fewer complex rounds" [2]. Threefish operates entirely on unsigned 64-bit integers and involves only three operations: rotation of k bits to the left, bitwise exclusive OR, and addition modulo 2^{64}. Therefore, the plaintext P and the cipher key K are converted to N_w 64-bit words. Note that the number of words N_w and the number of rounds N_r depend on the key size (Table I). The size of a plaintext block is given by $N_b = 8 \cdot N_w$ bytes.

The key schedule generates the subkeys from a block cipher key $K = (k_0, k_1, \ldots, k_{N_w-1})$ and a 128-bit tweak $T = (t_0, t_1)$. K and T are extended with one parity word (Algorithm 1, steps 1 and 2). Each subkey is a combination of N_w words of the extended key, two words of the extended tweak, and a counter s (Algorithm 1, steps 5 to 9). Note that the extended key and the extended tweak are rotated by one word position between two consecutive subkeys.

Algorithm 1 Key schedule of Threefish.

Input: A block cipher key $K = (k_0, k_1, \ldots, k_{N_w-1})$; a tweak $T = (t_0, t_1)$; the constant $C_{240} = \texttt{1BD11BDAA9FC1A22}$.

Output: $N_r/4 + 1$ subkeys $k_{s,0}, k_{s,1}, \ldots, k_{s,N_w-1}$, where $0 \leq s \leq N_r/4$.

1. $k_{N_w} \leftarrow C_{240} \oplus \bigoplus_{i=0}^{N_w-1} k_i$;

2. $t_2 \leftarrow t \oplus t_1$;

3. **for** $s \leftarrow 0$ **to** $N_r/4$ **do**

4.   **for** $i \leftarrow 0$ **to** $N_w - 4$ **do**

Fig. 1. One of the 72 encryption rounds of Threefish-256.

Fig. 2. One of the 72 decryption rounds of Threefish-256.

5. $k_{s,i} \leftarrow k_{(s+i)\bmod(N_w+1)}$;

6. **end for**

7. $k_{s,N_w-3} \leftarrow k_{(s+N_w-3)\bmod(N_w+1)} \boxplus t_{s\bmod 3}$;

8. $k_{s,N_w-2} \leftarrow k_{(s+N_w-2)\bmod(N_w+1)} \boxplus t_{(s+1)\bmod 3}$;

9. $k_{s,N_w-1} \leftarrow k_{(s+N_w-1)\bmod(N_w+1)} \boxplus s$;

10. **end for**

11. **return** $k_{s,0}, k_{s,1}, \ldots, k_{s,N_w-1}$, where $0 \le s \le N_r/4$;

A series of N_r rounds (Fig. 1 and Algorithm 2, steps 4 to 19) and a final subkey addition (Algorithm 2, step 21) are applied to produce the ciphertext. The core of a round is the simple non-linear mixing function $\mathrm{Mix}_{d,j}$ (Algorithm 2, steps 13 and 14). It consists of an addition, a rotation by a constant $R_{d\bmod 8,j}$ (repeated every eight rounds and defined in [2, Table 4]), and a bit-wise exclusive OR. A word permutation $\pi(i)$ (see [2, Table 3]) is then applied to obtain the output of the round (Algorithm 2, step 17). Furthermore, a subkey is injected every four rounds (Algorithm 2, step 7).

Algorithm 2 Encryption with the Threefish block cipher.

Input: A plaintext block $P = (p_0, p_1, \ldots, p_{N_w-1})$; $N_r/4 + 1$ subkeys $k_{s,0}, k_{s,1}, \ldots, k_{s,N_w-1}$, where $0 \le s \le N_r/4$; $4N_w$ rotation constants $R_{i,j}$, where $0 \le i \le 7$ and $0 \le j \le N_w/2$.

Output: A ciphertext block $C = (c_0, c_1, \ldots, c_{N_w-1})$.

1. **for** $i \leftarrow 0$ **to** $N_w - 1$ **do**

2. $v_{0,i} \leftarrow p_i$;

3. **end for**

4. **for** $d \leftarrow 0$ **to** $N_r - 1$ **do**

5.   **for** $i \leftarrow 0$ **to** $N_w - 1$ **do**

6.   **if** $d \bmod 4 = 0$ **then**

7. $e_{d,i} \leftarrow v_{d,i} \boxplus k_{d/4,i}$;   (Key injection)

8.   **else**

9. $e_{d,i} \leftarrow v_{d,i}$;   (Rename)

10.   **end if**

11.   **end for**

12.   **for** $j \leftarrow 0$ **to** $N_w/2 - 1$ **do**

13. $f_{d,2j} \leftarrow e_{d,2j} \boxplus e_{d,2j+1}$; $(\mathrm{Mix}_{d,j})$

14. $f_{d,2j+1} \leftarrow f_{d,2j} \oplus (e_{d,2j+1} \lll R_{d\bmod 8,j})$;

15.   **end for**

16.   **for** $i \leftarrow 0$ **to** $N_w - 1$ **do**

17. $v_{d+1,i} \leftarrow f_{d,\pi(i)}$;   (Permute)

18.   **end for**

19. **end for**

20. **for** $i \leftarrow 0$ **to** $N_w - 1$ **do**

21. $c_i \leftarrow v_{N_r,i} \boxplus k_{N_r/4,i}$;   (Key injection)

22. **end for**

23. **return** $C = (c_0, c_1, \ldots, c_{N_w-1})$;

Fig. 2 describes a decryption round of Threefish-256. It consists of the inverse word permutation followed by the inverse MIX functions. Note that subkeys are injected in reverse order.

III. THE SKEIN FAMILY OF HASH FUNCTIONS

The Unique Block Iteration (UBI) chaining mode allows one to build a compression function out of a tweakable encryption function $E(T, K, P)$. Let M be a message of arbitrary length up to $2^{99} - 8$ bits. If the number of bits in M is not a multiple of 8, we append a bit 1 followed by a (possibly empty) string of 0's. This step guarantees that M contains N_M bytes. Then, we pad M with p zero bytes so that $N_M + p$ is a multiple of the block size N_b. We can now split M into N_b-byte blocks $M_0, \ldots, M_{\beta-1}$, where $\beta = (N_M + p)/N_b$. Each block M_i is processed with a unique tweak value T_i encoding how many bytes have been processed so far, a type field (see [2] for details),

and two bits specifying whether it is the first and/or last block. The UBI chaining mode is computed as:

$$H_0 \leftarrow G,$$
$$H_{i+1} \leftarrow M_i \oplus E(H_i, T_i, M_i),$$

where G is a starting value of N_b bytes. In this work, we consider the normal hashing mode and refer the reader to [2] for a description of Skein-MAC and tree hashing with Skein. Skein is built on three invocations of UBI:

- Define a 32-byte configuration string C that contains the length of the digest size (in bits), a schema identifier, and a version number [2, Table 7]. Compute the N_b-byte block G_0 as $G_0 \leftarrow \mathrm{UBI}(0, C, T_{\mathrm{cfg}}2^{120})$. Note that G_0 only depends on the digest size and can easily be precomputed.
- The message is then processed as follows: $G_1 \leftarrow \mathrm{UBI}(G_0, M, T_{\mathrm{msg}}2^{120})$.
- A third call to UBI is required to achieve hashing-appropriate randomness: $H \leftarrow \mathrm{UBI}(G_1, 0, T_{\mathrm{out}}2^{120})$. This transform allows one to produce arbitrary digest sizes (up to 2^{64} bits). If a single output block H is not enough, one can use Threefish in counter mode to produce the digest.

IV. THE ChaCha STREAM CIPHER

The ChaCha family of stream ciphers was designed by Bernstein [3] to improve the diffusion per round of Salsa20 [6], while preserving the encryption rate. ChaCha operates on 32-bit words, and expands a 256-bit key $(k_0, \ldots, k_7)$ and a 64-bit nonce $(\mathrm{IV}_0, \mathrm{IV}_1)$ into a 2^{70}-byte stream. A b-byte message is then encrypted (or decrypted) by XORing it with the first b bytes of the stream.

ChaCha generates the stream by blocks of 64 bytes. In order to process the ith block, ChaCha acts on a 4×4 matrix M of 32-bit integers defined as follows:

$$\begin{pmatrix} m_0 & m_1 & m_2 & m_3 \\ m_4 & m_5 & m_6 & m_7 \\ m_8 & m_9 & m_{10} & m_{11} \\ m_{12} & m_{13} & m_{14} & m_{15} \end{pmatrix} = \begin{pmatrix} c_0 & c_1 & c_2 & c_3 \\ k_0 & k_1 & k_2 & k_3 \\ k_4 & k_5 & k_6 & k_7 \\ t_0 & t_1 & \mathrm{IV}_0 & \mathrm{IV}_1 \end{pmatrix},$$

where

- $c_0 = \text{61707865}$, $c_1 = \text{3320646E}$, $c_2 = \text{79622D32}$, and $c_3 = \text{6B206574}$ are predefined constants;
- $t = (t_0, t_1)$ is a 64-bit counter encoding the index i (i.e. $i = 2^{32}t_1 + t_0$).

ChaCha transforms the matrix M through a series of N_r rounds (Algorithm 3). The algorithm is based on a nonlinear operation called quarter-round function and described by Algorithm 4. Matrix M is copied into matrix V. Then, the even- and odd-numbered rounds of ChaCha apply the quarter-round function to each row and northwest-to-southeast diagonal of V, respectively. Eventually, a new block of the stream is generated by adding V to the original matrix M (Algorithm 3, step 15), and the block counter is incremented (Algorithm 3, steps 17 to 20).

Algorithm 3 Computation of a 64-byte block of the stream of ChaCha.

Input: A key, a nonce, and a block counter stored in a matrix M.

Output: A 64-byte block of the stream.

1. **for** $i \leftarrow 0$ **to** 15 **do**
2. $v[i] \leftarrow m[i]$;
3. **end for**
4. **for** $i \leftarrow 0$ **to** $N_r/2 - 1$ **do**
5.   QUARTERROUND(v_0, v_4, v_8, v_{12});
6.   QUARTERROUND(v_1, v_5, v_9, v_{13});
7.   QUARTERROUND$(v_2, v_6, v_{10}, v_{14})$;
8.   QUARTERROUND$(v_3, v_7, v_{11}, v_{15})$;
9.   QUARTERROUND$(v_0, v_5, v_{10}, v_{15})$;
10.   QUARTERROUND$(v_1, v_6, v_{11}, v_{12})$;
11.   QUARTERROUND(v_2, v_7, v_8, v_{13});
12.   QUARTERROUND(v_3, v_4, v_9, v_{14});
13. **end for**
14. **for** $i \leftarrow 0$ **to** 15 **do**
15. $v[i] \leftarrow v[i] \boxplus m[i]$;
16. **end for**
17. $m_{12} \leftarrow m_{12} \boxplus 1$;
18. **if** $m_{12} = 0$ **then**
19. $m_{13} \leftarrow m_{13} \boxplus 1$;
20. **end if**
21. **Return** M and V;

Algorithm 4 The ChaCha quarter-round function.

Input: Four 32-bit integers a, b, c, and d.

Output: QUARTERROUND(a, b, c, d).

1. $a \leftarrow a \boxplus b$;
2. $d \leftarrow (d \oplus a) \lll 16$;
3. $c \leftarrow c \boxplus d$;
4. $b \leftarrow (b \oplus c) \lll 12$;
5. $a \leftarrow a \boxplus b$;
6. $d \leftarrow (d \oplus a) \lll 8$;
7. $c \leftarrow c \boxplus d$;
8. $b \leftarrow (b \oplus c) \lll 7$;
9. **Return** a, b, c, and d;

TABLE II
PROPERTIES OF THE BLAKE HASH FUNCTIONS

Algorithm	Word size w [bits]	Message [bits]	Block size b [bits]	Digest size [bits]	Salt [bits]	# rounds N_r	Rotation distances			
							δ_0	δ_1	δ_2	δ_3
BLAKE-224	32	$< 2^{64}$	512	224	128	14	16	12	8	7
BLAKE-256	32	$< 2^{64}$	512	256	128	14	16	12	8	7
BLAKE-384	64	$< 2^{128}$	1024	384	256	16	32	25	16	11
BLAKE-512	64	$< 2^{128}$	1024	512	256	16	32	25	16	11

Bernstein proposed 8-, 12-, and 20-round variants of ChaCha. Aumasson *et al.* introduced a novel method for differential cryptanalysis of ChaCha and broke the 7-round variant [7]. Ishiguro *et al.* [8], [9] improved the attack and concluded that Salsa20 and ChaCha "are not presently under threat".

V. THE BLAKE FAMILY OF HASH FUNCTIONS

The BLAKE family combines three previously studied components, chosen by Aumasson *et al.* for their complementarity [1]: the iteration mode HAIFA, the internal structure of the hash function LAKE, and a modified version of Bernstein's stream cipher ChaCha as compression function. BLAKE is a family of four hash functions, namely BLAKE-224, BLAKE-256, BLAKE-384, and BLAKE-512 (Table II). The main differences lie in the length of words w, the number of rounds N_r, and in some constants involved in the algorithm. In the following, we denote by BLAKE-n the algorithm with an n-bit digest.

BLAKE-n involves only two arithmetic operations: the addition modulo 2^w of two w-bit unsigned integers and the bitwise exclusive OR of two w-bit words. The latter is sometimes followed by a rotation of δ_j bits to the right. The four possible rotation distances depend on the digest size and are defined in Table II. The compression function of BLAKE-n produces a new chain value $h' = h'_0, \ldots, h'_7$ from a message block $m = m_0, \ldots, m_{15}$, a chain value $h = h_0, \ldots, h_7$, a salt $s = s_0, \ldots, s_3$, a counter $t = t_0, t_1$, and 16 constants c_i defined in [1, p. 8]. This process consists of three steps. First, a 16-word internal state $v = v_0, \ldots, v_{15}$ is initialized as follows:

$$\begin{pmatrix} v_0 & v_1 & v_2 & v_3 \\ v_4 & v_5 & v_6 & v_7 \\ v_8 & v_9 & v_{10} & v_{11} \\ v_{12} & v_{13} & v_{14} & v_{15} \end{pmatrix} \leftarrow \begin{pmatrix} h_0 & h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 & h_7 \\ s_0 \oplus c_0 & s_1 \oplus c_1 & s_2 \oplus c_2 & s_3 \oplus c_3 \\ t_0 \oplus c_4 & t_0 \oplus c_5 & t_1 \oplus c_6 & t_1 \oplus c_7 \end{pmatrix}.$$

Then, a series of N_r rounds is performed. Each of them consists of a transformation of the internal state v based on the G_i function described by Algorithm 5, where σ_r denotes a permutation of $\{0, \ldots, 15\}$ parametrized by the round index r (see [1, Table 2.1]). A column step updates the four columns of matrix v as follows: $G_0(v_0, v_4, v_8, v_{12})$, $G_1(v_1, v_5, v_9, v_{13})$, $G_2(v_2, v_6, v_{10}, v_{14})$, and $G_3(v_3, v_7, v_{11}, v_{15})$. Note that each call to G_i updates a distinct column of matrix v. Since we focus on compact implementations of BLAKE in this work, we interleave the computation of G_0, G_1, G_2, and G_3. This approach allows us to design an ALU with four pipeline stages and to achieve high clock frequencies. Then, a diagonal step updates the four diagonals of v: $G_4(v_0, v_5, v_{10}, v_{15})$, $G_5(v_1, v_6, v_{11}, v_{12})$, $G_6(v_2, v_7, v_8, v_{13})$, and $G_7(v_3, v_4, v_9, v_{14})$. Here again, each call to G_i modifies

a distinct diagonal of the matrix, allowing us to interleave the computation of G_4, G_5, G_6, and G_7.

Algorithm 5 The G_i function.

Input: A function index i and four w-bit integers a, b, c, and d.

Output: $G_i(a, b, c, d)$.

1. $a \leftarrow a \boxplus b$;

2. $a \leftarrow a \boxplus (m_{\sigma_r(2i)} \oplus c_{\sigma_r(2i+1)})$;

3. $d \leftarrow (d \oplus a) \ggg \delta_0$;

4. $c \leftarrow c \boxplus d$;

5. $b \leftarrow (b \oplus c) \ggg \delta_1$;

6. $a \leftarrow a \boxplus b$;

7. $a \leftarrow a \boxplus (m_{\sigma_r(2i+1)} \oplus c_{\sigma_r(2i)})$;

8. $d \leftarrow (d \oplus a) \ggg \delta_2$;

9. $c \leftarrow c \boxplus d$;

10. $b \leftarrow (b \oplus c) \ggg \delta_3$;

At the end of the last round, a new chain value $h' = h'_0, \ldots, h'_7$ is computed from the internal state v and the previous chain value h (finalization step):

$$\begin{aligned} h'_0 &\leftarrow h_0 \oplus s_0 \oplus v_0 \oplus v_8, & h'_4 &\leftarrow h_4 \oplus s_0 \oplus v_4 \oplus v_{12}, \\ h'_1 &\leftarrow h_1 \oplus s_1 \oplus v_1 \oplus v_9, & h'_5 &\leftarrow h_5 \oplus s_1 \oplus v_5 \oplus v_{13}, \\ h'_2 &\leftarrow h_2 \oplus s_2 \oplus v_2 \oplus v_{10}, & h'_6 &\leftarrow h_6 \oplus s_2 \oplus v_6 \oplus v_{14}, \\ h'_3 &\leftarrow h_3 \oplus s_3 \oplus v_3 \oplus v_{11}, & h'_7 &\leftarrow h_7 \oplus s_3 \oplus v_7 \oplus v_{15}. \end{aligned}$$

In order to guarantee that the length ℓ of a message is a multiple of the block size b, Aumasson *et al.* define the following padding scheme [1]:

- append a bit 1 followed by a sufficient number of 0 bits such that the length is congruent to $b - 2w - 1$ modulo b;
- a padding bit followed by the $2w$-bit unsigned big-endian representation of ℓ is then added; in the case of BLAKE-256 and BLAKE-512, the padding bit is equal to 1; otherwise, it is set to 0.

The hash can now be computed iteratively: the padded message is divided into β 16-word blocks $m^{(0)}, \ldots, m^{(\beta-1)}$ and the chain value $h^{(0)}$ is set to the same initial value as SHA-n. The counter $t^{(i)}$ denotes the number of message bits in $m^{(0)}, \ldots, m^{(i)}$. If the last block contains only padding bits, then $t^{(\beta-1)}$ is set to zero. The message digest consists of the n least significant bits of the output $h^{(\beta)}$.
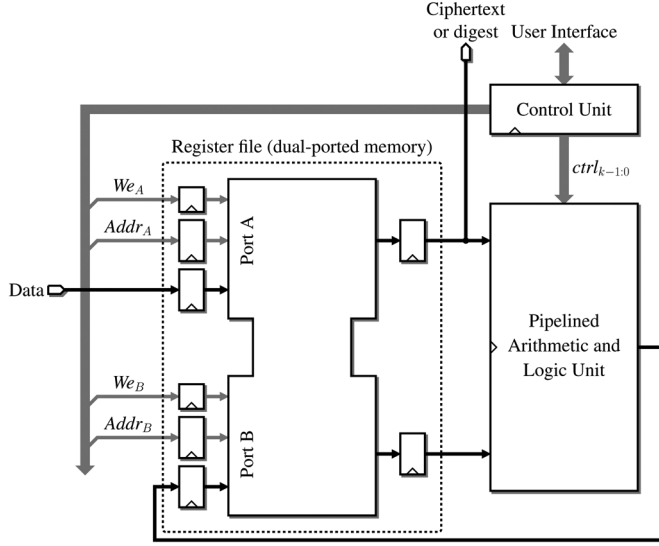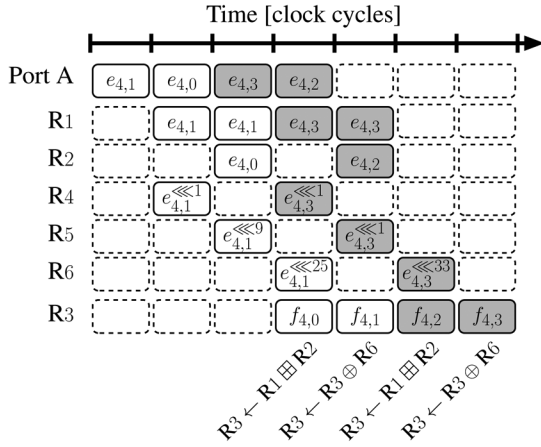
Fig. 3. General architecture of our coprocessors.

Fig. 4. Computation of $\text{Mix}_{4,0}$ and $\text{Mix}_{4,1}$ (Threefish-256). Reprinted from [4].

VI. HARDWARE IMPLEMENTATION

All of our architectures consist of a register file organized into w-bit words and implemented by means of dual-ported memory, an ALU, and a control unit (Fig. 3). The user loads messages, plaintext blocks or ciphertext blocks into port A. A few control bits allows her to select the algorithm and the desired level of security. When the coprocessors are hashing or encrypting a message, the intermediates results are always written to port B. The result is stored the register file and can be read word by word on port A. In the following, we assume that our coprocessors are provided with padded messages. A hardware wrapper interface for BLAKE, Skein, and several other hash functions comprising communication and padding is described in [10].

A. Arithmetic and Logic Units for Threefish and Skein

Our first ALU, originally described in [4], implements Threefish encryption and Skein. In the following, $\text{R}i$ denotes a 64-bit register. Fig. 4 illustrates our scheduling of the two mixing functions $\text{Mix}_{4,0}$ and $\text{Mix}_{4,1}$ of the fifth round of Threefish-256:

- The operand $e_{4,1}$ is loaded in register R1; at the same time, we start the computation of $e_{4,1} \lll R_{4,0}$; this opera-
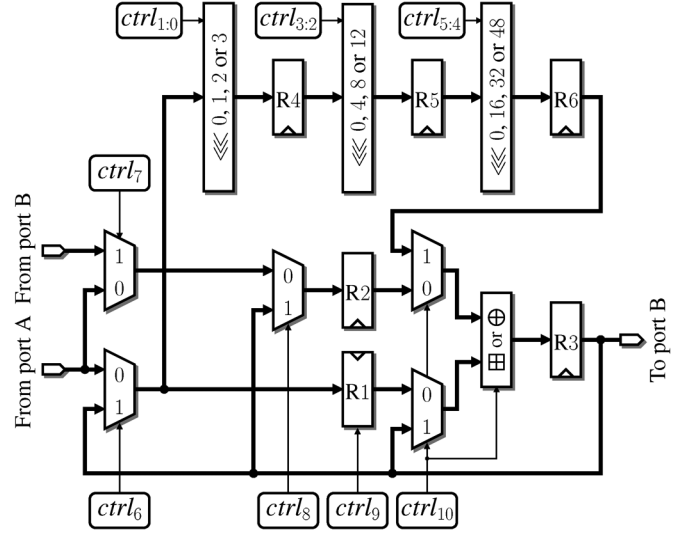

Fig. 5. Arithmetic and logic unit for Threefish encryption. Reprinted from [4].

tion requires three clock cycles and intermediate results are stored in R4, R5, and R6.
- Then, $e_{4,0}$ is loaded in register R2; the content of R1 is not modified (i.e. R1 must be controlled by an enable signal).
- We execute the instruction $\text{R3} \leftarrow \text{R1} \boxplus \text{R2}$ and obtain $f_{4,0}$.
- R3 and R6 contain $f_{4,0}$ and $e_{4,1} \lll R_{4,0}$, respectively. The instruction $\text{R3} \leftarrow \text{R3} \oplus \text{R6}$ allows us to compute $f_{4,1}$.

We schedule $\text{Mix}_{4,1}$ as soon as $e_{4,0}$ has been read, and manage to keep the pipeline continuously busy. In summary, our ALU must be able to carry out any rotation of a 64-bit word and to perform the following operation (Fig. 5):

$$\text{R3} \leftarrow \begin{cases} \text{R1} \boxplus \text{R2} & \text{when } ctrl_{10} = 0, \\ \text{R3} \oplus \text{R6} & \text{otherwise,} \end{cases} \quad (1)$$

where $ctrl_{10}$ denotes a control bit. Let us define two 64-bit operands a and b such that:

$$(a, b) = \begin{cases} (\text{R1}, \text{R2}) & \text{when } ctrl_{10} = 0, \\ (\text{R3}, \text{R6}) & \text{otherwise.} \end{cases}$$

It is well-known that $a \boxplus b = (a \vee b) \boxplus (a \wedge b)$ and $a \oplus b = (a \vee b) \boxminus (a \wedge b)$ [11]. Thus, (1) can be rewritten as follows:

$$\text{R3} \leftarrow (a \vee b) \boxplus ((a \wedge b) \oplus ctrl_{10}) \boxplus ctrl_{10}. \quad (2)$$

Fig. 6 describes the implementation of (2) on a Virtex-6 device. Since there is a single control signal to choose the arithmetic operation and to select a and b, (2) involves only five variables, and is advantageously implemented by 64 LUT6_2 primitives and dedicated carry logic.

In order to reduce the number of operands stored in the register file, we interleave the key schedule (Algorithm 1) and the encryption process (Algorithm 2). This approach allows us to generate the subkeys on-the-fly. It is however necessary to compute t_2 and k_{N_w} before the first key injection. The easiest way to compute t_2 would be to load t_0 and t_1 in registers R1 and R2, respectively, and to execute the instruction $\text{R3} \leftarrow \text{R1} \oplus \text{R2}$. Unfortunately, this solution requires one more control bit to select the inputs of the arithmetic operator, and it is not possible to implement the multiplexers and the adder on the same LUT6_2 primitive anymore. Since the critical path of our co-

Fig. 6. Computation of $R3 \leftarrow R1 \boxplus R2$ or $R3 \leftarrow R3 \oplus R6$ on a Virtex-6 device. Reprinted from [4].

processor is located in the 64-bit adder, an extra level of LUTs would decrease the clock frequency. However, we are able to compute t_2 using only the functionalities defined by (1). Since $t_2 = (t_0 \boxplus 0) \oplus (t_1 \lll 0)$, it suffices to execute the following instructions:

$$R4 \leftarrow t_1 \lll 0,$$
$$R1 \leftarrow t_0, \quad R2 \leftarrow 0, \quad R5 \leftarrow R4 \lll 0,$$
$$R3 \leftarrow R1 \boxplus R2, \quad R6 \leftarrow R5 \lll 0,$$
$$R3 \leftarrow R3 \oplus R6.$$

This approach assumes that we can read simultaneously two values from the register file. Thanks to the multiplexer controlled by $ctrl_7$, we can load data from port A or port B into register R_2 (Fig. 5). A similar strategy allows us to compute k_{N_w}.

The implementation of the key injection is more straightforward. Note that the multiplexers controlled by $ctrl_6$ and $ctrl_8$ allow us to bypass the register file and to use the content of R3 as an input to the ALU. Let us consider for instance the first key injection of Threefish-256: $e_{0,2}$ is defined as $p_2 \boxplus k_{0,2} = p_2 \boxplus k_2 \boxplus t_1$ and is computed as follows:

$$R1 \leftarrow k_2, \quad R2 \leftarrow t_1,$$
$$R3 \leftarrow R1 \boxplus R2$$
$$R1 \leftarrow R3, \quad R2 \leftarrow p_2,$$
$$R3 \leftarrow R1 \boxplus R2.$$

Fig. 7 describes how we schedule the instructions of Threefish-256 decryption.

The UBI chaining mode can be combined with the final key injection of Threefish encryption. It suffices to modify step 21 of Algorithm 2 as follows:

$$e_{N_r,i} \leftarrow v_{N_r,i} \boxplus k_{\frac{N_r}{4},i};$$
$$c_i \leftarrow e_{N_r,i} \oplus p_i.$$

The only difference between this operation and the mixing function $\mathrm{MIX}_{d,j}$ is that no permutation is applied to the second operand of the bitwise exclusive OR.

The inverse of the MIX function is purely sequential. Therefore, Threefish decryption has less parallelism than encryption, and it is not possible to compute an addition modulo 2^{64} and a rotation in parallel. We suggest to modify our ALU as follows to fully support both encryption and decryption (Fig. 8):

- The inverse of the Mix function and the inverse of the key injection require a subtraction modulo 2^{64}. Our modified ALU is able to perform a new operation: $R3 \leftarrow R1 \boxminus R2$. An additional control bit $ctrl_{16}$ allows us to add R2 or its two's complement to R1. It is therefore not possible to implement our arithmetic operator by means of 64 LUT6_2 anymore, and the slice count and the critical path are expected to increase.
- The output of the inverse Mix function is provided either by the arithmetic operator (e.g. $e_{4,0}$ on Fig. 2) or the rotation unit (e.g. $e_{4,1}$ on Fig. 2). The multiplexer controlled by $ctrl_{17}$ allows us to select the word we store in the register file.
- Since the inverse of the Mix function is sequential, we have to perform the rotation in a single clock cycle. We suggest to take advantage of the SRL16E primitive available on Xilinx devices to implement a FIFO whose depth is dynamically adjusted according to the algorithm selected by the user: one and three stages for decryption and encryption, respectively.

B. Arithmetic and Logic Units for BLAKE and ChaCha

Let us consider the G_i function of BLAKE-n to define the instruction set of our coprocessors. Since we focus on compact coprocessors for the BLAKE family in this article, we perform a single step of Algorithm 5 at each clock cycle. We will show later that the input operand b is already stored in an internal register of our ALU when we start the computation of $G_i(a,b,c,d)$. Therefore, each operation involves the result of the previous one, and our ALU will include a feedback mechanism to bypass the register file of the coprocessor.

Assume that the w-bit word computed by the ALU is stored in register R5, and denote by RF_A and RF_B the operands provided by the register file. From the data flow diagram of Algorithm 5, we easily identify three operations (Fig. 9):

1) Save the content of R5 in the register file and compute $R5 \leftarrow R5 \boxplus \mathrm{RF}_A$.
2) Compute $R5 \leftarrow R5 \boxplus (\mathrm{RF}_A \oplus \mathrm{RF}_B)$.
3) Save the content of R5 in the register file and compute $R5 \leftarrow (R5 \oplus \mathrm{RF}_A) \ggg \delta_j$.

Recall now that the four calls to G_i in a column step or a diagonal step can be computed in parallel. In order to keep the critical path as short as possible, we suggest to design an ALU with four pipeline stages and to interleave the computation of four G_i functions (Fig. 10). The heart of the ALU is the arithmetic operator performing the addition or the bitwise XOR of two w-bit words described in Section VI-A. Our operator computes:

$$R3 \leftarrow \begin{cases} R1 \boxplus R2 & \text{when } ctrl_2 = 0, \\ R1 \oplus R2 & \text{otherwise} \end{cases}$$
$$= (R1 \vee R2) \boxplus ((R1 \wedge R2) \oplus ctrl_2) \boxplus ctrl_2,$$

where

- R1 stores the data provided by the register file. Since a flip-flop is always associated with a LUT, we can perform
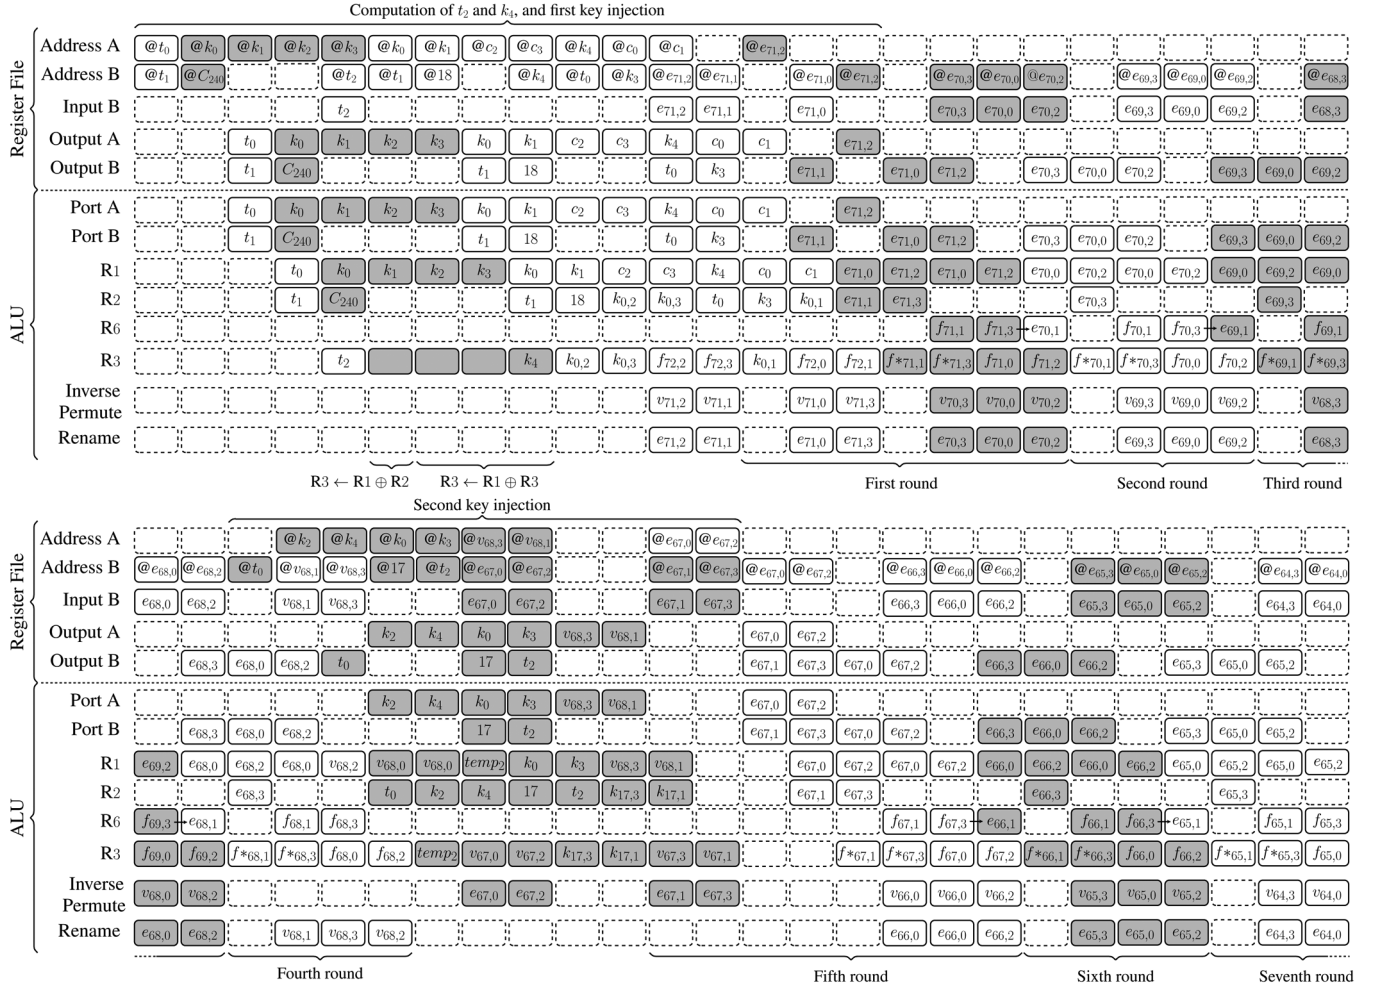
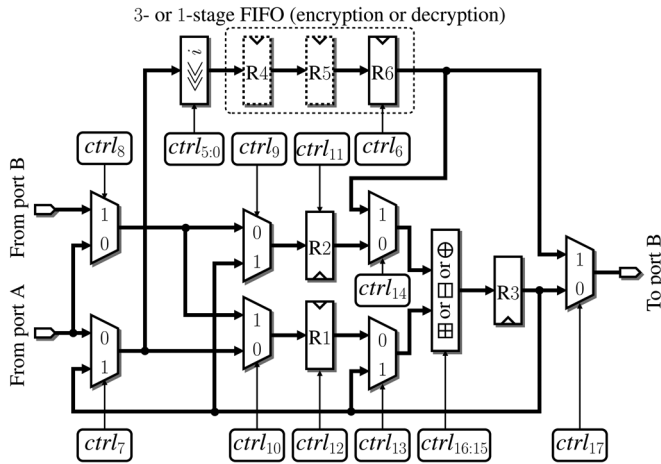Fig. 7. Scheduling of Threefish-256 decryption.

Fig. 8. Arithmetic and logic unit for Threefish encryption and decryption.

some simple pre-processing without increasing the number of slices of the ALU: a control bit $ctrl_0$ selects either RF_A or $\mathrm{RF}_A \oplus \mathrm{RF}_B$. This allows us to compute $m_{\sigma_r(2i)} \oplus c_{\sigma_r(2i+1)}$ and $m_{\sigma_r(2i+1)} \oplus c_{\sigma_r(2i)}$ for free (Algorithm 5, steps 2 and 7).

- R2 almost always stores the result of a previous operation. However, we have to disable the feedback mechanism during the initialization step: the computation of

$v_8 \leftarrow s_0 \oplus c_0$ involves for instance only two words stored in the register file. An array of AND gates controlled by $ctrl_1$ allows us to force the second operand to zero in such cases.

If needed, the content of register R3 is then rotated to the left in two steps. Our implementation is based on the following observation:

$$\mathrm{R3} \ggg \delta_i = (\mathrm{R3} \ggg (\delta_i - \delta_3)) \ggg \delta_3,$$

where $0 \leq i \leq 3$. At first glance, this design choice may look awkward. However, it will allow us to easily build a unified processor for the BLAKE family. The key point is that the content of R3 is copied into R5 when the three control bits $ctrl_{5:3}$ are equal to 0 (Fig. 10).

Note that the pipeline has three possible configurations, denoted by ①, ②, and ③ in Fig. 10:

① In order to minimize the area of our ALU, we can insert a w-bit register after the first stage of the rotation. Since the latter involves w LUTs, there is no hardware overhead on a Virtex-6 device.

② The addition modulo 2^w can be computed in two clock cycles. Let $a = a_{\mathrm{low}} + 2^{w/2} a_{\mathrm{high}}$ and $b = b_{\mathrm{low}} + 2^{w/2} b_{\mathrm{high}}$. We store a_{high} and b_{high} in two $w/2$-bit registers, and compute a sum word s_{low} and a carry bit such that $2^{w/2} c + s_{\mathrm{low}} = a_{\mathrm{low}} + b_{\mathrm{low}}$. A flip-flop and a $w/2$-bit register store
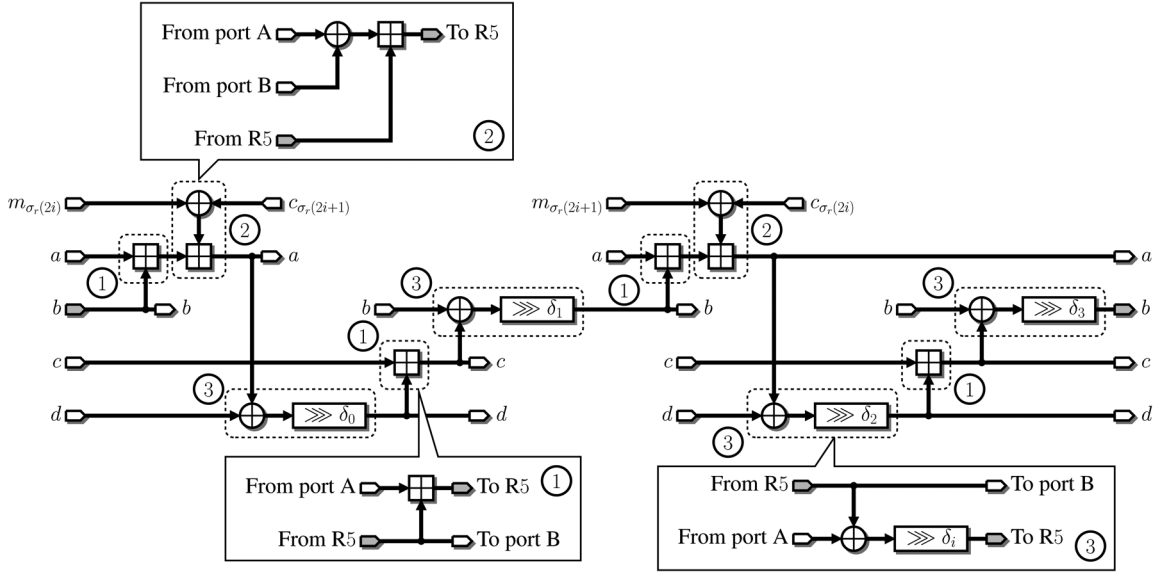
Fig. 9. Implementation of the G_i function of BLAKE-n by means of three instructions. R5 denotes an internal register of the ALU.
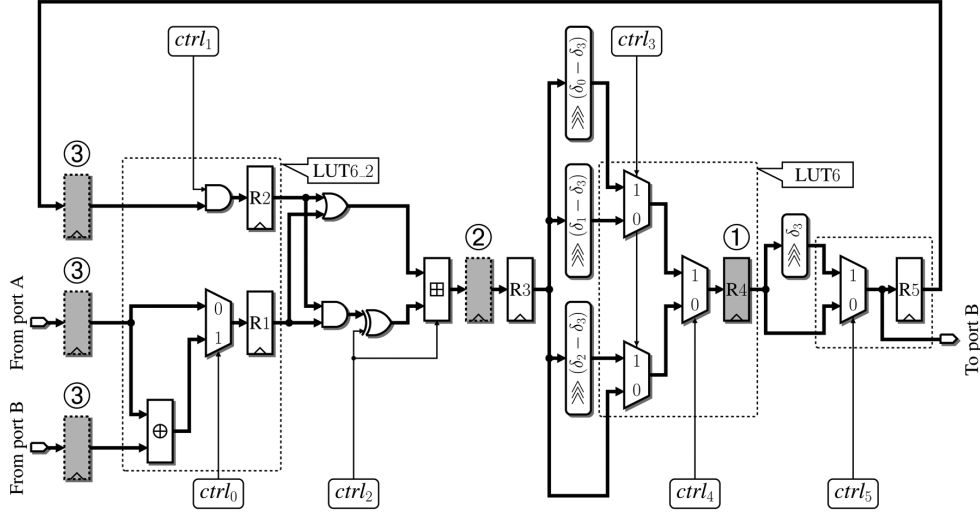

Fig. 10. Arithmetic and Logic Unit for BLAKE-n.

c and s_{low}, respectively. The most significant bits of the sum are then given by $s_{\text{high}} = a_{\text{high}} + b_{\text{high}} + c$. This approach allows us to reduce the worst-case carry path at the price of three $w/2$-bit registers and a flip-flop.

③ Routing a signal from a memory block to a slice is sometimes expensive in terms of wire delay. If the critical path is located between the register file and register R1, this pipeline configuration will help boosting the clock frequency. The output data path of a Virtex-6 memory block has an optional internal pipeline register. Therefore, the only hardware overhead is the w-bit register between R5 and the array of AND gates controlled by $ctrl_1$.

In order to avoid pipeline bubbles between column and diagonal steps, it suffices to process the four calls to G_i of the diagonal step in the following order [5]: G_7, G_4, G_5, and G_6. We check for instance that the ALU outputs the new value of v_4 (last instruction of G_0) at time $\tau + 3$. If we load v_3 from the register file, we can start the computation of G_7 at time $\tau + 4$. We easily check that this scheduling also avoids pipeline bubbles

between a diagonal step and a column step (Fig. 11). Since each call to G_i involves ten instructions, we need 80 clock cycles to perform a round of BLAKE-n.

Our first architecture can be modified to support the four algorithms of the BLAKE family (Fig. 12). The 64-bit datapath is built out of two 32-bit datapaths, thus allowing us to perform a single 64-bit operation or two 32-bit operations at each clock cycle. The mode of operation is selected according to an additional control bit $ctrl_6$, the latter being provided by the user. The ALU includes two 32-bit adders. Let a_{low}, a_{high}, b_{low}, b_{high}, s_{low}, and s_{high} denote unsigned 32-bit integers. When the user chooses BLAKE-224 or BLAKE-256 ($ctrl_6 = 0$), two messages are processed in parallel and the ALU performs two 32-bit additions:

$$s_{\text{low}} \leftarrow a_{\text{low}} + b_{\text{low}} + ctrl_2,$$
$$s_{\text{high}} \leftarrow a_{\text{high}} + b_{\text{high}} + ctrl_2.$$

When the coprocessor executes BLAKE-384 or BLAKE-512 ($ctrl_6 = 1$), the ALU carries out a 64-bit addition. The first
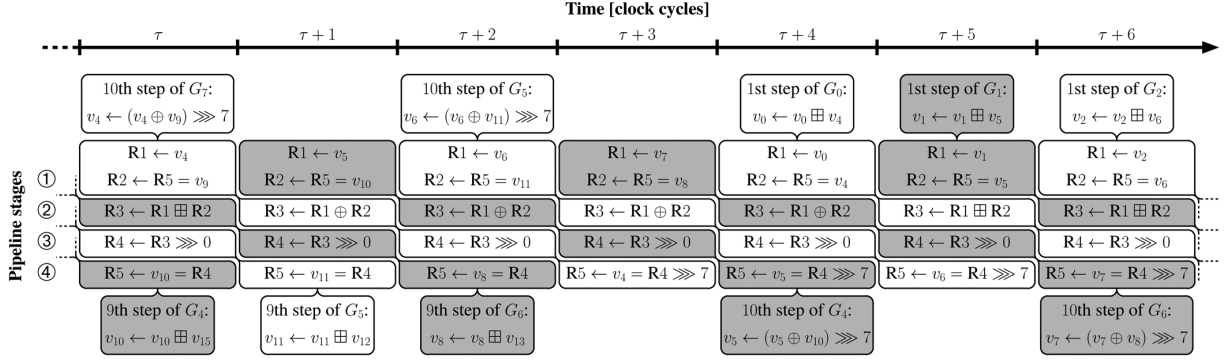
Fig. 11. Avoiding pipeline bubbles between a diagonal step and a column step.
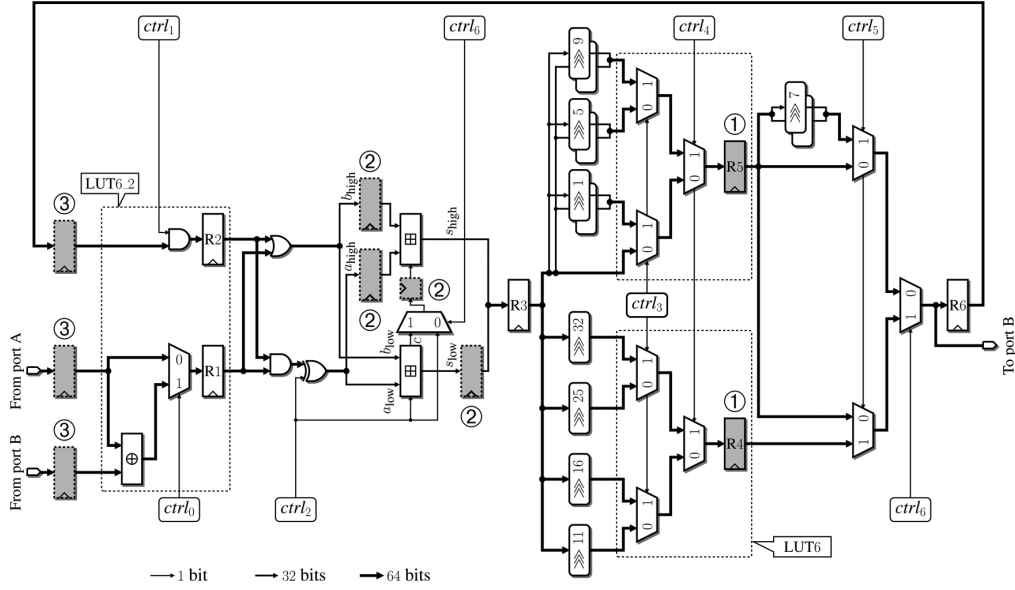

Fig. 12. Unified arithmetic and logic unit for the BLAKE family.

adder generates the least significant bits of the sum and a carry bit c such that:

$$2^{32} \cdot c + s_{\text{low}} = a_{\text{low}} + b_{\text{low}} + ctrl_2.$$

The second adder computes the most significant bits of the sum:

$$s_{\text{high}} \leftarrow a_{\text{high}} + b_{\text{high}} + c.$$

We use the rotation unit of our first processor to deal with BLAKE-224 and BLAKE-256. Note that the content of R3 is always copied into R6 when $ctrl_{5:3} = (000)_2$. Thus, we share this datapath between all algorithms of the BLAKE family, and need only 64 LUTs to implement the rotation unit of BLAKE-384 and BLAKE-512. When $ctrl_5$ is equal to one, $ctrl_{4:3}$ encodes the index i of the rotation distance δ_i (Table III). Consequently, we can use the same instruction flow for all algorithms and select the width of the datapath according to $ctrl_6$. Note that the three pipeline configurations defined for our first coprocessor are also available here.

The QUARTERROUND function of ChaCha requires only two of the instructions we defined for the G_i function. Thus, the design of a ChaCha coprocessor is rather straightforward (Fig. 13). Since it is not necessary to compute $\text{RF}_A \oplus \text{RF}_B$ anymore, the ALU has a single 32-bit input. The only difficulty

TABLE III
ROTATION DISTANCES OF THE UNIFIED BLAKE COPROCESSOR

$ctrl_{5:3}$	Rot. dist.	BLAKE-224/256	BLAKE-384/512
$(000)_2$	0	$R_6 \leftarrow R_3$ (common datapath)	
$(100)_2$	δ_0	$R_6 \leftarrow R_3 \ggg 7$	$R_6 \leftarrow R_3 \ggg 11$
$(101)_2$	δ_1	$R_6 \leftarrow R_3 \ggg 8$	$R_6 \leftarrow R_3 \ggg 16$
$(110)_2$	δ_2	$R_6 \leftarrow R_3 \ggg 12$	$R_6 \leftarrow R_3 \ggg 25$
$(111)_2$	δ_3	$R_6 \leftarrow R_3 \ggg 16$	$R_6 \leftarrow R_3 \ggg 32$

is to increment the 64-bit counter $2^{32}m_{13} + m_{12}$ (Algorithm 3, steps 17 to 20). Assume that the constant 1 is stored in the register file. The control bit $ctrl_0$ allows us to disable the feedback mechanism and to load the constant 0 in register R2. Execute the following instructions:

$$\text{R1} \leftarrow 1 \quad \text{R2} \leftarrow 0,$$
$$\text{R3} \leftarrow \text{R1} \oplus \text{R2},$$
$$\text{R4} \leftarrow \text{R3},$$
$$\text{R5} \leftarrow \text{R4},$$
$$\text{R1} \leftarrow m_{12} \quad \text{R2} \leftarrow \text{R5}.$$

Registers R1 and R2 store m_{12} and the constant 1, respectively. Note that the output carry of the 32-bit adder can now be stored in a flip-flop F. Furthermore, when $ctrl_2$ is set to one, our ALU
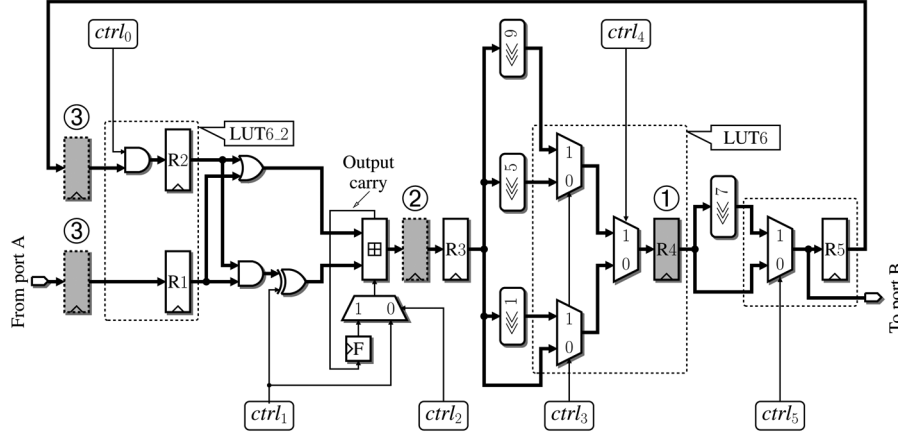
Fig. 13. Arithmetic and Logic Unit for ChaCha.

Fig. 14. Unified arithmetic and logic unit for the BLAKE and ChaCha families.

performs an "add with carry" instruction. We can now compute $m_{12} + 1$, save the output carry in F, and increment m_{13} if necessary:

$$(F, R3) \leftarrow R1 \boxplus R2 \quad R1 \leftarrow m_{13} \quad R2 \leftarrow 0,$$
$$R3 \leftarrow R1 \boxplus R2 \boxplus F \quad R4 \leftarrow R3,$$
$$\text{Register file} \leftarrow R4(m_{12}) \quad R4 \leftarrow R3,$$
$$\text{Register file} \leftarrow R4(m_{13}).$$

Three pipeline configurations are again available. The second one needs specific attention: since the adder is pipelined, the computation of $m_{12} \boxplus 1$ requires two clock cycles. It is therefore mandatory to introduce a NOP before loading m_{13} into register R1.

It is of course possible to build a unified coprocessor for ChaCha and the BLAKE family (Fig. 14). A new control bit $ctrl_7$ allows the user to select the mode of operation of the ALU: encryption or hashing. Since the coprocessor has a 64-bit datapath to support BLAKE-384 and BLAKE-512, it is possible to encrypt two messages in parallel with ChaCha.

C. Register Files and Control Units

We will consider our unified coprocessor for the BLAKE and ChaCha algorithms to describe how we design our control units. The same approach can easily be applied to the other coprocessors considered in this work. Virtex-6 FPGAs embed several configurable memory blocks that can for instance store 1024
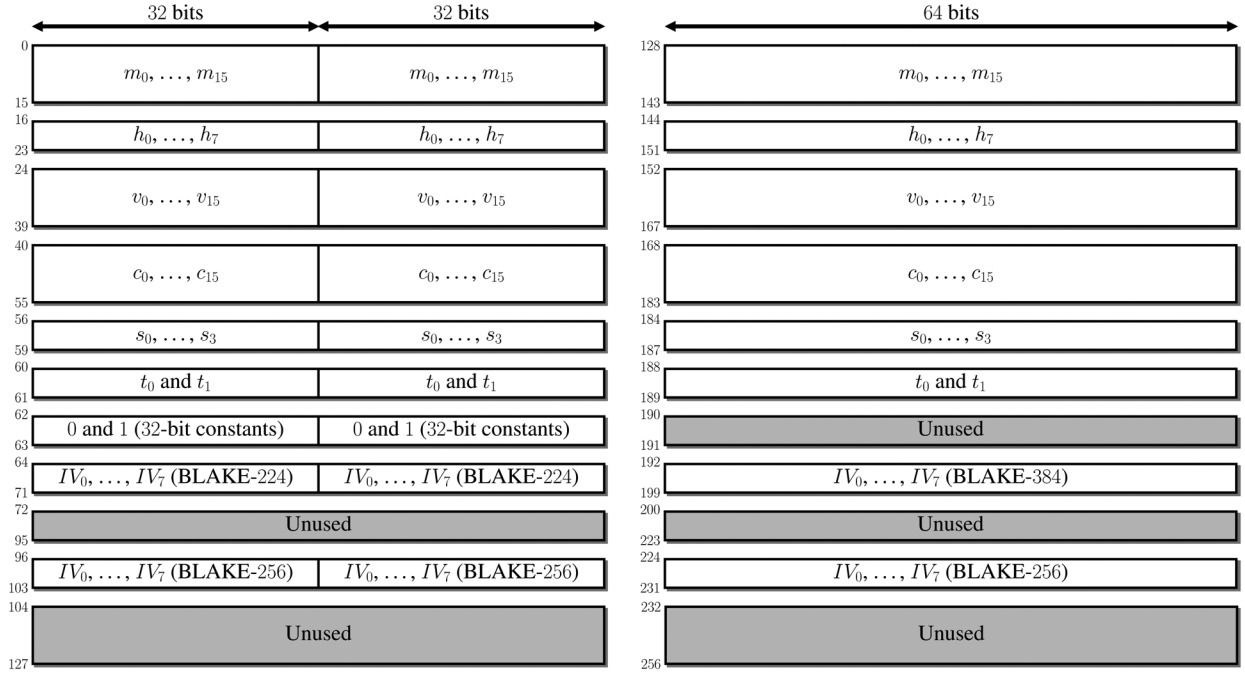
Fig. 15. Register file of the unified coprocessor for the BLAKE and ChaCha families.


Fig. 16. From 26-to 20-bit instructions. Shaded cells denote control bits provided by the user. (a) Address and control bits of our unified coprocessor for BLAKE and ChaCha; (b) Address and control bits provided by the control unit.

TABLE IV
NUMBER OF INSTRUCTIONS OF THE ALGORITHMS OF THE BLAKE AND CHACHA FAMILIES

Algorithm	# instructions
BLAKE-224/256	1184
BLAKE-384/512	1344
8-round ChaCha	311
12-round ChaCha	439
20-round ChaCha	695

36-bit words or 2048 18-bit words. Our control unit mainly consists of a program counter that addresses an instruction memory implemented by means of a memory block.

A straightforward way to deal with the permutations involved in the BLAKE family is to unroll the round loop. Table IV summarizes the number of instructions required by the algorithms supported by our coprocessor if we follow this approach. Note that it suffices to store the code of BLAKE-384/512 and 20-round ChaCha (2039 instructions): a simple finite-state machine allows us to jump to the finalization step when the desired number of rounds has been performed. The main challenge is therefore to define control words of at most 18 bits in order to implement our instruction memory by means of a single memory block. A clever organization of the register file (Fig. 15) and a simple compression algorithm allows us to achieve this goal: it suffices to list all the possible values taken

by the 7-bit word $We_B\|ctrl_5\|\ldots\|ctrl_0$ for each pipeline configuration. Since the number of choices is always smaller then 32, it is possible to label each pattern with a 5-bit string. Two blocks of dual-ported memory configured as 256 entries of 32 bits store the message, the chaining value, the constants, and all the intermediate variables of BLAKE and ChaCha. Thus, our coprocessor requires 26 control bits (Fig. 16(a)):

- 8 address bits and a write enable signal for port A of the register file;
- 8 address bits and a write enable signal for port B of the register file;
- 8 control bits for the ALU.

Two control bits are provided by the user: $ctrl_7$ allows her to select between BLAKE and ChaCha, and $ctrl_6$ specifies the configuration of the datapath (2×32 bits or 64 bits). Our organization of the data in the register file enables us to define a 20-bit instruction:

- The most significant address bit depends on the algorithm being executed, and is therefore provided by the user.
- We use ports A and B to load new data (message, salt, and counter) and save the intermediate variables computed by the ALU, respectively. Consequently, the write enable signal of port A is also given by the user.
- Let us denote by $a_{7:0}$ the eight address bits of ports A. Note that a_6 is equal to one only when we read an initial vector and assume that the digest size is selected according to an
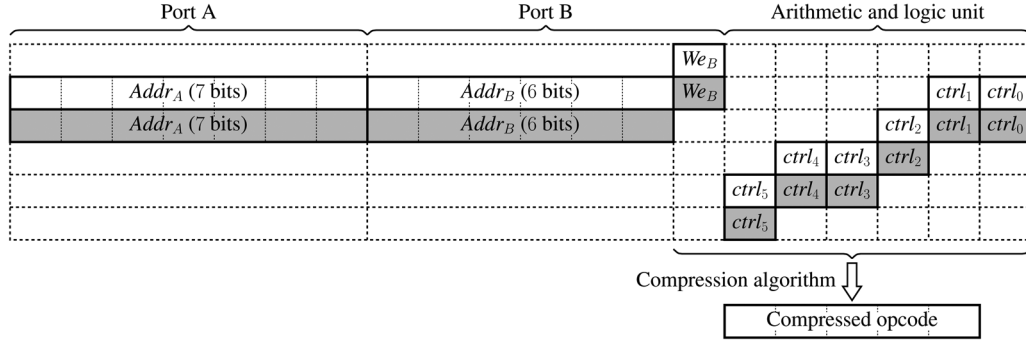
Fig. 17. Generation of the compressed instruction memory.

TABLE V
PLACE-AND-ROUTE RESULTS FOR OUR THREEFISH AND SKEIN COPROCESSORS ON A VIRTEX-6 FPGA (XC6VLX75T-2). ALL DESIGNS REQUIRE THREE MEMORY BLOCKS

Supported Algorithms	Area [Slices]	Freq. [MHz]	Throughput [Mbits/s]							
			Skein-256-256	Skein-512-512	Threefish-256		Threefish-512		Threefish-1024	
					Enc.	Dec.	Enc.	Dec.	Enc.	Dec.
Threefish encryption	145	294	–	–	153	–	175	–	160	–
Threefish	277	267	–	–	139	145	158	125	145	116
Skein and Threefish encryption	150	295	75	85	154	–	175	–	161	–
Skein and Threefish	292	279	70	80	145	152	166	130	152	121

additional control bit $ctrl_8$. The address bit a_5 is computed as follows:

$$a_5 \leftarrow \begin{cases} a_5 & \text{when } a_6 = 0, \\ ctrl_8 & \text{otherwise.} \end{cases}$$

Thanks to this simple mechanism, the instruction flow does not depend on the digest size. Initial vectors are always read from port A.

- Since the initial vectors are neither modified nor read from port B, the second most significant address bit is always equal to zero.

Consequently, we can store 20-bit words in the instruction memory (Fig. 16(b)). We designed a simple compression algorithm to encode the write enable signal of port B and the six control bits $ctrl_{5:0}$ by means of five bits. A C program generates the content of the instruction memory and the VHDL description of the decompression circuit. The latter involves only seven 5-input LUTs, and stores the control bits of the ALU and the write enable signal of port B in a register. Because of this pipeline stage, it is necessary to generate the write enable signal one clock cycle in advance when we have to store a word in the register file. Our C program takes this parameter into account and organizes the control bits in the instruction memory according to the pipeline configuration. Then, it generates the compressed instruction memory. Fig. 17 describes the instruction flow for the first pipeline configuration of our coprocessor:

- As explained above, the write enable signal is generated one clock cycle in advance to take the internal pipeline stage of the decompression unit into account.
- All inputs of the register file are registered, and the two control bits $ctrl_0$ and $ctrl_1$ must therefore be generated one clock cycle after the addresses. We take advantage of the latency of our decompression unit to synchronize the control signals.

We followed the same approach to build our control units for Threefish and Skein. The register file is organized into 64-bit words, and stores a plaintext block, an internal state ($e_{d,i}$, where $0 \leq i \leq N_w - 1$), an extended block cipher key, an extended tweak, the constant C_{240}, and all possible values of s involved in the key schedule. Thanks to this approach, the word permutation $\pi(i)$ and the word rotation of the key schedule are conveniently implemented by addressing the register file accordingly. Since the round constants repeat every eight rounds (Algorithm 2, step 14), we decided to unroll eight iterations of the main loop of Threefish (Algorithm 2, steps 4 to 19). The rotation constants $R_{d,i}$ are included in the microcode executed by the control unit. Note that our register file is designed for Threefish-1024 (i.e. $N_w = 16$ and $N_r = 80$). It is therefore straightforward to implement the two other variants of the algorithm on our architecture.

VII. RESULTS AND COMPARISONS

We captured our architecture in the VHDL language and prototyped our coprocessors on a Xilinx Virtex-6 FPGA with average speedgrade. Tables V and VI summarize our place-and-route results measured with ISE 14.2. Note that we considered the least favorable case, where the message consists of a single block, to compute the throughput of Skein.

Most of the architectures described in the open literature focus on a single level of security (Table VII). We took advantage of the intrinsic parallelism of BLAKE to interleave the computation of four instances of the G_i function. Thanks to this approach, we designed an ALU with four pipeline stages and achieved higher clock speeds than the coprocessors listed in Table VII. A careful scheduling allowed us to totally avoid pipeline bubbles and memory collisions. We also addressed FPGA-specific issues and described how to share slices between addition and bitwise exclusive OR of two operands. We followed the same strategy to design our coprocessors

Supported algorithms	Pipeline config.	Area [slices]	# block RAMs	Freq. [MHz]	Throughput [Mbits/s]				
					BLAKE-224 and BLAKE-256	BLAKE-384 and BLAKE-512	8-round ChaCha	12-round ChaCha	20-round ChaCha
ChaCha	①	49	2	362	–	–	595	422	266
	②	77	2	316	–	–	520	368	232
	③	77	2	345	–	–	569	403	254
BLAKE-224 and BLAKE-256	①	47	2	338	146	–	–	–	–
	②	49	2	341	147	–	–	–	–
	③	50	2	349	150	–	–	–	–
BLAKE-384 and BLAKE-512	①	79	3	331	–	252	–	–	–
	②	91	3	331	–	252	–	–	–
	③	91	3	329	–	250	–	–	–
BLAKE (all levels of security)	①	94	3	312	2×134	237	–	–	–
	②	126	3	332	2×143	252	–	–	–
	③	129	3	343	2×148	261	–	–	–
BLAKE and ChaCha (all levels of security)	①	144	3	335	2×144	255	2×551	2×390	2×246
	②	156	3	289	2×124	220	2×475	2×337	2×212
	③	168	3	304	2×131	231	2×500	2×354	2×223

|	Supported algorithm(s)	FPGA	Area [slices]	36k memory blocks	Frequency [MHz]	Throughput [Mbits/s]
Latif *et al.* [12][†]	Skein-256-256	xc5vlx110-3	821	Not specified	119	1610
Jungk [13][†,‡]	Skein-512-256	xc5v	555	–	271	237
Jungk [14][†,‡]	Skein-512-256	xc6v	406	–	318	277
Kaps *et al.* [15]	Skein-512-256	xc6vlx75t-1	207	1	166	17
Kaps *et al.* [15]	Skein-512-256	xc6vlx75t-1	193	–	193	21
Kerckhof *et al.* [16][†]	Skein-512-512	xc6vlx75t-1	240	–	160	179
Aumasson *et al.* [1]	BLAKE-256	xc5vlx110	390	–	91	412
Jungk [14]	BLAKE-256	xc6v	235	–	231	518
Jungk [14]	BLAKE-256	xc6v	404	–	185	823
Kaps *et al.* [15]	BLAKE-256	xc6vlx75t-1	163	1	197	327
Kaps *et al.* [15]	BLAKE-256	xc6vlx75t-1	166	–	268	445
Aumasson *et al.* [1]	BLAKE-512	xc5vlx110	939	–	59	468
Kerckhof *et al.* [16]	BLAKE-512	xc6vlx75t-1	192	–	240	183

Fig. 18. Compact implementations of several cryptographic hash functions on Virtex-6 FPGAs (512-bit digests).

for Threefish and Skein. As a consequence, our coprocessors provide the end-user with hashing and encryption at all levels of security, while offering a better area-time trade-off.

We report in Fig. 18 the latest lightweight implementation results of several cryptographic hash functions. Besides our coprocessors for BLAKE-512 and Skein-512-512, we selected

Grøstl [17], JH [16], SHA-2-512 [18], and SHA-3-512 (Keccak [$r = 1024$, $c = 576$]) [19]. In this context, BLAKE is obviously the best choice for lightweight implementations on FPGA. Since our unified architecture for the BLAKE family (Fig. 12) requires less than 100 Virtex-6 slices, BLAKE is also an excellent candidate for cryptographic coprocessors supporting several levels of security. We already proposed lightweight implementations of ECHO & AES [20] and Grøstl & AES [17]. According to our results, the unified coprocessor for BLAKE and ChaCha offers the best area-time trade-off.

VIII. CONCLUSION

The stream cipher ChaCha, the block cipher Threefish, and the hash functions BLAKE and Skein are based on the same arithmetic operations. In this work, we showed that the same design philosophy allows one to design lightweight coprocessors for hashing and encryption. The key element of our approach is to take advantage of the parallelism of the algorithms to deeply pipeline the ALU to achieve a high clock frequency, and avoid data dependencies by interleaving independent tasks. Furthermore, we described how to design compact control units thanks to a careful organization of the register file, loop unrolling, and

a simple compression algorithm. Our architectures are mainly designed for embedded systems. Thus, it would be interesting to conduct side-channel and fault injection attacks in future work.

REFERENCES

[1] J.-P. Aumasson, L. Henzen, W. Meier, and R. Phan, SHA-3 Proposal BLAKE (version 1.4) Jan. 2011 [Online]. Available: http://www.131002.net/blake

[2] N. Ferguson, S. Lucks, B. Schneier, D. Whiting, M. Bellare, T. Kohno, J. Callas, and J. Walker, The Skein Hash Function Family (version 1.3) Oct. 2010 [Online]. Available: http://www.skein-hash.info

[3] D. Bernstein, ChaCha, a Variant of Salsa20 Jan. 2008 [Online]. Available: http://cr.yp.to/papers.html#chacha

[4] N. At, J.-L. Beuchat, and İ. San, "Compact implementation of Threefish and Skein on FPGA," in *Proc. 5th IFIP Int. Conf. New Technologies, Mobility and Security—NTMS 2012*, A. Levi, M. Badra, M. Cesana, M. Ghassemian, O. Gürbüz, N. Jabeur, M. Klonowski, A. Mana, S. Sargento, and S. Zeadally, Eds. New York, NY, USA: IEEE eXpress Conference Publishing, 2012.

[5] J.-L. Beuchat, E. Okamoto, and T. Yamazaki, "Compact implementations of BLAKE-32 and BLAKE-64 on FPGA," in *Proc. 2010 Int. Conf. Field-Programmable Technology—FPT 2010*, J. Bian, Q. Zhou, and K. Zhao, Eds. Piscataway, NJ, USA: IEEE Press, 2010, pp. 170–177.

[6] D. Bernstein, The Salsa20 family of stream ciphers Dec. 2007 [Online]. Available: http://cr.yp.to/snuffle/salsafamily-20071225.pdf

[7] J.-P. Aumasson, S. Fischer, S. Khazaei, W. Meier, and C. Rechberger, "New features of Latin dances: Analysis of Salsa, ChaCha, and Rumba," in *Fast Software Encryption—FSE 2008*, K. Nyberg, Ed. New York, NY, USA: Springer, 2008, vol. 5086, Lecture Notes in Computer Science, pp. 470–488.

[8] T. Ishiguro, S. Kiyomoto, and Y. Miyake, "Latin dances revisited: New analytic results of Salsa20 and ChaCha," in *Information and Communications Security—ICICS 2011*, S. Qing, W. Susilo, G. Wang, and D. Liu, Eds. New York, NY, USA: Springer, 2011, vol. 7043, Lecture Notes in Computer Science, pp. 255–266.

[9] T. Ishiguro, Modified Version of "Latin Dances Revisited: New Analytic Results of Salsa20 and ChaCha" Report 2012/065, 2012, cryptology ePrint Archive.

[10] B. Baldwin, A. Byrne, L. Lu, M. Hamilton, N. Hanley, M. O'Neill, and W. Marnane, A Hardware Wrapper for the SHA-3 Hash Algorithms Report 2010/124, 2010, cryptology ePrint Archive.

[11] H. Warren, *Hacker's Delight*. New York, NY, USA: Addison-Wesley, 2002.

[12] K. Latif, M. Tariq, A. Aziz, and A. Mahboob, "Efficient hardware implementation of secure hash algorithm (SHA-3) finalist-Skein," in *Proc. Int. Conf. Computer, Communication, Control Automation—3CA2011*, 2011.

[13] B. Jungk, "Compact implementations of Grøstl, JH and Skein for FPGAs," in *Proc. ECRYPT II Hash Workshop*, 2011 [Online]. Available: http://www.ecrypt.eu.org/hash2011/program.shtml

[14] B. Jungk, "Evaluation of compact FPGA implementations for all SHA-3 finalists," in *Proc. 3rd SHA-3 Candidate Conf.*, Mar. 2012.

[15] J.-P. Kaps, P. Yalla, K. Surapathi, B. Habib, S. Vadlamudi, and S. Gurung, "Lightweight implementations of SHA-3 finalists on FPGAs," in *Proc. 3rd SHA-3 Candidate Conf.*, Mar. 2012 [Online]. Available: http://csrc.nist.gov/groups/ST/hash/sha-3/Round3/March2012/Program_SHA3_March2012.html

[16] S. Kerckhof, F. Durvaux, N. Veyrat-Charvillon, F. Regazzoni, G. Meurice de Dormale, and F.-X. Standaert, "Compact FPGA implementations of the five SHA-3 finalists," in *Proc. ECRYPT II Hash Workshop*, 2011 [Online]. Available: http://www.ecrypt.eu.org/hash2011/program.shtml

[17] N. At, J.-L. Beuchat, E. Okamoto, İ San, and T. Yamazaki, A Low-Area Unified Hardware Architecture for the AES and the Cryptographic Hash Function Grøstl Report 2012/535, 2012, cryptology ePrint Archive.

[18] H. Technology, FULL DATASHEET—Tiny Hash Core Family for Xilinx FPGA revision 2.0 (11/06/2010).

[19] İ. San and N. At, "Compact Keccak hardware architecture for data integrity and authentication on FPGAs," *Inf. Security J.: A Global Perspective*, vol. 21, no. 5, pp. 231–242, 2012.

[20] J.-L. Beuchat, E. Okamoto, and T. Yamazaki, "A low-area unified hardware architecture for the AES and the cryptographic hash function ECHO," *J. Cryptographic Eng.*, vol. 1, no. 2, pp. 101–121, 2011.

Nuray At received the B.Sc. degree in electrical and electronics engineering from Anadolu University, Eskişehir, Turkey, the M.Sc. degree in electrical and computer engineering from the University of California, Davis, CA, USA, in 1998, the M.Sc. degree from the Georgia Institute of Technology, GA, USA, in 2001, and the Ph.D. degree in electrical and electronics engineering from Anadolu University, Eskişehir, Turkey, in 2005.

She held a visiting researcher position at the Selmer Center, University of Bergen. She is currently a faculty member in Electrical and Electronics Engineering Department at Anadolu University, Eskişehir, Turkey. Her research interests include communication systems, information security, and signal processing.

Jean-Luc Beuchat received the M.Sc. and Ph.D. degrees in computer science from the Swiss Federal Institute of Technology, Lausanne, Switzerland, in 1997 and 2001, respectively.

He is a member of the Security Division at ELCA Informatique SA. His current research interests include computer arithmetic and cryptography.

Dr. Beuchat is an Associate Editor of the IEEE TRANSACTIONS ON COMPUTERS.

Eiji Okamoto received the B.S., M.S. and Ph.D degrees in electronics engineering from the Tokyo Institute of Technology, Tokyo, Japan, in 1973, 1975 and 1978, respectively.

He worked and studied communication theory and cryptography for NEC central research laboratories since 1978. In 1991 he became a professor at Japan Advanced Institute of Science and Technology, then at Toho University. Now he is a professor at Graduate School of Systems and Information Engineering, University of Tsukuba. His research interests are cryptography and information security.

Dr. Okamoto is a member of IEEE and a Coeditor-In-Chief of *International Journal of Information Security*.

İsmail San received the B.Sc. degree in electrical and electronics engineering from Anadolu University, Turkey, in 2008, where he is currently working toward the Ph.D. degree in the Electrical Engineering Department.

His main interests are in architecture design for information and communication security with special focus on efficient implementations for cryptographic applications.

Teppei Yamazaki received the MCE degree from the University of Tsukuba, Japan, in 2012.

He currently works at Trend Micro corp. and studies information security including dynamic malware analysis and intrusion detection systems.